

EVIDENCE & IMPLEMENTATION / 01

The demonstrated result

[A finite subscription domain.](#) [Real model reasoning.](#) [Explicit local authority.](#)

Accepted implementation V11, 13 September 2026. This appendix is a readable evidence gateway for the committed result, not a new runtime execution.

Testflix is a finite subscription lifecycle executed through the common Hedgehog OS machinery. Real Gemini outputs select between supplied subscription plans. The selected result is consumed by typed work and by independent decisions for the user, bank, provider and device. A mock payment receipt becomes an input to entitlement issuance; the entitlement and current device grant become inputs to playback.

The main story covers purchase, an informational request, stop and fresh-session behavior, expiry, and explicit renewal. Separate controlled P04 branches exercise actual public E recomputation when a quote changes. A later live A/B changes the preference while preserving the other request fields, obtains a different plan, and consumes that result through payment and playback. Captured execution reproduces the A/B on fresh Hosts without calling the model.

The final owner execution applied 51 paths to main: 31 modifications and 20 additions. It retained 896 other tracked files, producing 947 tracked files. The historical V09 branch remains at 268f22a7940068fd58e00e6f96dc2d13a7d71890. The owner phase is PUSHED, with no blocker. The agreed implementation/admission scope was independently accepted.

Evidence: [owner state](#), [actual A/B](#), [main result](#).

EVIDENCE & IMPLEMENTATION / 02

Four reasoning duties, four independent Roots

The live semantic provider is Gemini 2.5 Flash. Four duties use the configured model; this is not a claim of four different model families. Each receives an ordered projection of its request and prior contributions.

Semantic duty	Consumed information	Returned contribution
Intent interpreter	Explicit preference, currency and ceiling	Priorities and any missing required evidence
Provider terms analyst	Request, finite catalogue and intent	Plan facts and budget eligibility
Client plan selector	Intent and analysed terms	A selected plan with decision factors
Contract reviewer	The same contract and all preceding contributions	Support or a blocking disagreement

A model output does not contain payment authority. Local code binds each output to its request, projection, role and transport record. The fourth duty can disagree: the preserved A/B refusal did so. A disagreement is not overwritten with a successful controlled response.

The reasoning duties are distinct from the four Root decisions:

Root	Identifier	Responsibility shown here
UserRoot	root:testflix:user	Accept the user's purchase selection
BankRoot	root:testflix:bank	Review the payment against current terms and consent
ProviderRoot	root:testflix:provider	Authorize entitlement issuance using the payment receipt
DeviceRoot	root:testflix:device	Authorize playback within session and device constraints

The actual A/B consumption record contains four accepted Root decisions with different decision IDs and target Root IDs. The same four Hosts remain throughout that branch's purchase, information, stop and fresh-session history. An independent captured run creates fresh Hosts.

Evidence: [semantic checkpoint](#), [consumption](#). Source: [semantic_roles_v01.py](#), [lifecycle_v01.py](#).

EVIDENCE & IMPLEMENTATION / 03

Purchase, reuse and explicit renewal

All stored amounts are integer minor currency units: 500 means EUR 5.00, not EUR 500. The catalogue contains three plans: EUR 5.00 at 720p without advertising, EUR 6.00 at 1080p with advertising, and EUR 9.00 at 2160p without advertising. Under the initial EUR 6.50 ceiling, the third is ineligible.

Stage	Observed behavior	Payment implication
P01 - purchase	AD_FREE selects plan:cd79, EUR 5.00, 720p, no advertising	One main mock payment
P02 - information	Admissible stored results answer an informational request	No new provider call or business effect
P03 - session lifecycle	Stop and a fresh session use the existing paid period and current device authority	No second purchase
P05 - expiry	At and after the entitlement boundary, the old right cannot authorize a new start	Zero effect in the exact-boundary controls
P06 - renewal	New explicit EUR 7.00 consent permits a new payment and paid period	Main total becomes two
After renewal	Stop and fresh-session behavior preserve the four Hosts and paid-period history	Main total stays two

Eight real model calls contribute to the successful main history: four for the purchase and four for renewal. The cumulative development journal is a different count, discussed below. No actual bank, streaming provider or physical television was operated.

The expiry controls evaluate a request at the boundary and immediately after it, with `executor_delta=0` and reason `current_entitlement`. Historical replay still verifies the old record. This distinguishes a historical fact from a currently usable permission.

Evidence: [main result and phase durations](#), [exact boundary](#), [after expiry](#), [post-renewal control](#).

EVIDENCE & IMPLEMENTATION / 04

When a quote changes

P04 is explicitly separate controlled public E evidence. It is not presented as another live LLM call inside the main narrative. A genuine quote-bound packet exists at EUR 5.00. Later source observations show changed terms while the pending packet still exists. The old packet is refused before a new business effect.

Public E computes the affected set and selectively recomputes the affected branch. Its result is consumed by a fresh BankRoot review. The previously paid period is preserved. The EUR 7.00 branch fails the EUR 6.50 consent limit. An independent EUR 6.00 branch receives a new review and packet, then produces one mock payment. The unchanged EUR 5.00 branch is the positive current-terms control.

Current quote	Existing consent	Result
EUR 5.00, unchanged	EUR 6.50	Independent positive branch executes
EUR 6.00, changed	EUR 6.50	Recompute, review and one independent mock payment
EUR 7.00, changed	EUR 6.50	Old packet refused; recomputed terms exceed consent

The EUR 6.00 control payment is not a third main-history payment. The later successful QUALITY A/B is also an independent branch with its own mock payment.

The evidence includes actual supplied D and E validations for the two completed recomputations. Coherent negative controls cover a missing dependency edge, an unreported change to preserved data, altered retained-work time/Host revision bindings, and a changed consumed result binding. These controls matter because a correctly formatted record alone is insufficient: it must describe the same work, source and review relation.

Evidence: [EUR 7.00 selected-field extract](#), [EUR 6.00 selected-field extract](#), [independent payment](#), [supplied D](#), [supplied E](#). Full report identities are in the evidence map.

EVIDENCE & IMPLEMENTATION / 05

The A/B result and the preserved disagreement

The controlled catalogue and EUR 6.50 ceiling stay the same. Changing preference from AD_FREE to QUALITY changes the chosen plan from plan:cd79 to plan:ab31: higher resolution, advertising present, price EUR 6.00. The consumed quote result, payment amount, entitlement plan, session quality and playback all reflect the new choice.

The first QUALITY attempt did not simply succeed. The interpreter associated QUALITY with both high resolution and an ad-free experience. The selector favoured resolution; the reviewer objected. Its actual response included:

[The upstream selection of 'plan:ab31' does not fully implement the 'QUALITY' preference because it includes ads.](#)

This is an excerpt from [response 13](#). It belongs to the original, pre-clarification input contract. The historical replay verifies that refusal under that original contract.

The shared QUALITY contract was subsequently made explicit: resolution descending, then absence of advertising, then price, then plan ID. Advertising became a tie-breaker rather than a veto on higher resolution. The validators continued to reject disagreement. Responses 14 and 15 had already completed interpretation and terms analysis under the clarified input. V10 validated and reused those responses, then obtained only the missing selector and reviewer as attempts 19 and 20.

The journal contains 20 attempts, 17 responses and three transport failures. Attempts 16 and 17 were recorded Google 504 failures; attempt 18 was a ReadTimeout. These are cumulative project-run counts, not the cost of a single successful purchase. The two V10 calls succeeded, and all original attempt/response/failure bodies remain preserved.

Evidence: [response 19](#), [response 20](#), [local raw-record and historical-refusal checks](#), [actual consumption](#).

EVIDENCE & IMPLEMENTATION / 06

Captures, replay and exact input identity

A capture binds the request, ordered role projection, raw output, response identity and transport metadata. The model is not asked to manufacture those local hashes. Saving an answer does not grant permission to execute it against a different current situation.

Two replay concepts must remain distinct:

Operation	What it does	Model/effect behavior
Captured reexecution	Executes the handler again on fresh Hosts using validated saved role responses	Zero model calls; fresh declared mock effects
Offline history replay	Verifies a sealed historical package and its expected manifest hash	Zero provider and capability calls; no new authority inferred

The V10 captured A/B reexecution reproduced the complete A/B history except for exactly two provenance fields: `purchase.generation_mode` and `purchase.semantics.mode`. No other fields were removed from comparison. Different provenance legitimately changes the seal hash.

The V11 saved-projection repair additionally distinguishes canonical JSON values such as integer 650 and float 650.0, or boolean false and integer 0. Python value equality alone does not preserve that distinction. Altered records are rejected before consumption, while the 12 genuine saved responses from main purchase, renewal and A/B still pass with identical canonical outputs.

Evidence: [captured comparison](#), [captured result](#), [offline replay](#), [V11 capture controls](#).

Timing and acceptance without double counting

Recorded operation	Duration	Measurement scope
V09 first purchase handler	169.10 s	Live P01 handler, before separate seal
V09 complete main story	1,174.80 s	Purchase through renewal, later sessions, sealing and replay
V09 P04 EUR 7.00 public E	317.99 s	Controlled selective recomputation
V09 P04 EUR 6.00 public E	318.91 s	Independent controlled selective recomputation
V10 actual A/B	290.806 s	Command wall time, including imports/finalizers
V10 captured A/B	276.216 s	Command wall time, including imports/finalizers

These are observed executions, not a production latency promise. The roughly five-minute A/B figures do not describe the entire purchase-to-renewal story. Model transport and local validation have different costs. The showcase build performs neither new model calls nor a new performance benchmark.

Evidence counts also describe distinct selections:

- V10: 36 focused tests and 108 setup/call/teardown phases.
- V11 preparation: 82 tests and 246 phases, plus eight disposable process-recovery controls.
- Actual owner postcommit: 49 unique tests and 147 passing phases.

Do not add these into a fictional number of unique system tests. V09 main/P04 and V10 A/B retain their recorded source-bound execution provenance. V11's owner checks establish final landing and selected current-boundary validation; they are not a fresh run of every previous D/E or Living suite.

Evidence: [V11 result](#), [owner node list](#), [owner complete phase log](#).

EVIDENCE & IMPLEMENTATION / 08

Scope, source map and review route

The reusable architectural contribution is the separation of proposal, Root decision, current execution eligibility and historical evidence across a concrete changing task. This demonstration shows that the common machinery can support this subscription domain alongside the earlier airline and supplier demonstrations. It does not establish that every future domain can be introduced without new domain contracts, integrations or engineering.

Concern	Accepted source
Finite request and plan contracts	hedgehog/domains/testflix/contracts_v01.py
Role projections and output binding	hedgehog/domains/testflix/semantic_roles_v01.py
Actual/captured provider transport	hedgehog/domains/testflix/live_semantic_adapter_v01.py
Four-Root lifecycle and dispatch	hedgehog/domains/testflix/lifecycle_v01.py
Story, repricing consumption and history	hedgehog/domains/testflix/evidence_v01.py
Declared business mocks	hedgehog/domains/testflix/mock_world_v01.py
Shared Host	hedgehog/work_execution_host_v01.py
Common typed work	hedgehog/kernel/work_composition_v01.py
D retained work	hedgehog/kernel/fractal_runtime_v02.py
E selective recomputation	hedgehog/kernel/continuous_delta_runtime_v01.py

Browse [the accepted domain](#) or the [accepted preflight and scope record](#).

The catalogue and clock are controlled. The clock is a local trusted test source, not an external attestation. All business effects are mocks; no real subscription or financial transaction is claimed. A new informational summary in the renewed period remains unsupported and outside this acceptance. No production or security certification is asserted.

A practical review order is: presentation / actual consumption / captured comparison / one preserved refusal / P04 selected-field extracts / owner phase log / accepted source. For a full reproduction audit, resolve the larger original reports using their archive/member/hash entries. The compact showcase is a readable evidence gateway, not a replacement for every original runtime object.

SUPPLEMENT / CAPTURE JOURNAL

Every attempt stays visible.

The complete journal spans the main preparation and the independent A/B completion. A returned answer can still be unusable or contain a refusal; a response count is not an execution-success count. All entries name gemini-2.5-flash.

#	Branch / duty	Recorded outcome	Seconds
1	main / Intent	Response	3.14
2	main / Intent	Response	3.68
3	main / Terms	Response	8.54
4	main / Selector	Response	7.18
5	main / Reviewer	Response	4.59
6	main / Intent	Response	2.12
7	main / Terms	Response	6.04
8	main / Selector	Response	4.37
9	main / Reviewer	Response	5.47
10	AB / Intent	Response	5.57
11	AB / Terms	Response	8.32
12	AB / Selector	Response	6.76
13	AB / Reviewer	Response: disagreement	11.94
14	AB / Intent	Response: reused in V10	2.41
15	AB / Terms	Response: reused in V10	11.40
16	AB / Selector	Google 504	n/a
17	AB / Selector	Google 504	n/a
18	AB / Selector	ReadTimeout	n/a
19	AB / Selector	Response: new in V10	6.59
20	AB / Reviewer	Response: new in V10	7.93

20 attempts = 17 responses + 3 transport failures. Per-call seconds are recorded response transport durations, not complete handler latency. The successful main story consumed eight model calls. The completed A/B consumed responses 14, 15, 19 and 20.

[SUPPLEMENT / IMMUTABLE REFERENCES](#)

Start with the source.

Every source link in this appendix is pinned to the accepted implementation. Evidence-map identities preserve the original archive and member of each copied or referenced record. A hash authenticates consistency with those originals; it does not attest an external-world event.

[Accepted commit](#)

e42d37fa98dfec7110b8cf75b1aceaa614f461be

[Accepted parent](#)

54e32dbcc0e4d68431ec2b9428eac965f88ee47c

[Accepted tree](#)

8f795d00b04b2604bb9d5c0fd973adaa98eb70d0

[Historical V09 commit](#)

268f22a7940068fd58e00e6f96dc2d13a7d71890

[V09 source archive](#)

RADIOLARIA_TESTFLIX_T3_REPAIR_AND_DEMONSTRATION_V09_RETURN_20260912T203950Z.tar.gz

SHA-256: f9c2d4519a404163fbfd3bdcef4632f9bba60025b3ad0e7ba982b85162937ec0

[V10 source archive](#)

RADIOLARIA_TESTFLIX_V10_AB_AND_CAPTURE_COMPLETION_RETURN_20260913T074750Z.tar.gz

SHA-256: e25b768177a59ed8ef15ff53c7de53b334ecfd0bb87bca7c8d024cf1f5a4b8af

[V11 source archive](#)

RADIOLARIA_TESTFLIX_V11_COMPLETED_ADMISSION_AND_OWNER_RETURN_20260913T094054Z.tar.gz

SHA-256: 0633a93be020766bb05fc1eeb121508675e45aefc3f6c107b5fe9b81cfc920c5

[OWNER source archive](#)

OWNER_EXECUTION_1789296250378385000.tar.gz

SHA-256: ce98c98dda6dbe82f485019ade7743cc4a327010f43dcedcb7c86ed5d9f0e783

[Review route](#)

Open evidence/evidence_map.json, follow actual A/B consumption, compare the captured record, read the original refusal and inspect the P04 extracts. Resolve full histories through the archive/member/hash references for a reproduction audit. The owner node list and phase log establish the final selected postcommit checks.

No model calls, financial effects or project runtime benchmarks were run to create this PDF. All prose is English. Conceptual artwork is separate from technical evidence.